

Fast Polyhedra Analysis with deep Q-networks

Jakub Kotal

Bachelor Thesis
November 2018

Supervisor:

Phd. Singh Gagandeep, Prof. Dr. Martin Vechev

Abstract

We use deep Q-networks to speed up numerical analysis. The key is reusing an existing link between program analysis and reinforcement learning. Expanding on these methods by testing and modifying the existing features and reward functions. Training with these new algorithms and learning a new and more optimal decision policy.

Contents

List of Figures	v
List of Tables	vii
1. Introduction	1
1.1. Problem Statement	1
1.2. Goals	2
1.3. Structure of this Document	2
2. Reinforcement Learning	3
2.1. Concepts	3
2.2. Q-function	4
2.3. Q-function approximation	4
2.3.1. Q-learning	5
2.3.2. Deep Q-networks	5
3. Polyhedra Analysis	7
3.1. Polyhedra representation	7
3.1.1. Constraint representation	8
3.1.2. Generator representation	8
3.2. Polyhedra domain	8
3.3. Polyhedra Decomposition	10
3.3.1. Decomposition operators	10
3.4. Reinforcement Learning for polyhedra analysis	11
3.4.1. Adapting polyhedra analysis for Reinforcement Learning	11

3.4.2. Existing methods	11
4. Fast polyhedra analysis with deep Q-networks	13
4.1. Incorporating Neural Networks inside Elina	13
4.2. Basic algorithm	13
4.2.1. Features	14
4.3. Actions	14
4.4. Reward	15
4.4.1. Experience replay memory	16
4.4.2. Separating target from max Q estimators	17
4.5. Further optimisations	18
4.5.1. Separating the problem into two	19
4.5.2. Feature selection	21
4.5.3. Reward Function modelling	23
4.5.4. Action selection algorithms	24
4.6. Neural network characteristics	28
4.7. Training	28
5. Results	31
5.1. Reward function selection	31
5.2. Action selection algorithm	32
5.3. Final algorithm	33
6. Conclusion and Future Work	35
A. Appendix	37
Bibliography	39

List of Figures

5.1. Reward one	32
5.2. Reward two	32
5.3. Reward three	32

List of Figures

List of Tables

Introduction

1

The number of domains controlled by computers has grown exponentially in the recent years. As the complexity of the structures surrounding us grows, the need for their automatision grows with it. And so, more and more the systems surrounding us, become controlled by computers. From simple things such as automatic doors to vastly more complex and important systems such as self-driving cars, medical software or nuclear weapons software. In some of these applications, a program bug could only be a slight nuisance, but in others its existence could have major implications. A few infamous examples of these exist where buffer overflows, rounding errors or division by zeros have caused space craft crashes or lethal radiation overdoses. As such, the importance of some of these devices demands their invulnerability and its verifiability. Unfortunately, as the complexity of these programs increases the need for complete and formal verification methods becomes critical.

Static analysis is a sub-branch of computer science. Its task is to analyse a program without its execution, its goal is to discover weaknesses inside of the code that could lead to vulnerabilities. As such, it should help with debugging and provide a better notion of safety about the code. Static analysis has seen a growing commercial use in the recent years mostly in safety-critical domains. Recent advances in this field use clever mathematical properties in order to increase the capabilities of such methods. Other techniques, leverage the domain of artificial intelligence to increase the fields capacities. For, until ai is capable of writing invulnerable software for us, we can at least use it verify ours.

1.1. Problem Statement

The main focus of static analysis is observing the effects that program expressions and statements have on the variables inside of the program. In order for these methods to work, a numerical abstract domain is needed that can capture the relationships between the variables. An important attribute of a domain is its expressivity, that is to say, how complex are the constraints between variables that the domain can represent. The more expressive a domain is the more complicated relationships it can represent. The most expressive domain is the polyhedra domain, representing the constraints with different polyhedra. However, its expressivity comes at a cost. It is notoriously time and space complex, having worst case exponential complexities in both. Other domains also exists that do not suffer from these problems. Unfortunately, to achieve this, they loose their expressivity. Modifications to the polyhedra domain are also possible, some of these modifications try to leverage precision loss against performance gain. However, finding the right balance between these two is not an easy task.

1.2. Goals

The goal of this thesis is to use some of the recent developments in new areas of reinforcement learning. To adapt and optimise these novel methods in order to use them with polyhedra analysis. The goal is then to optimise these methods with some problem specific knowledge from polyhedra analysis, in order to render the analysis more efficient and outperform other methods.

1.3. Structure of this Document

Chapter 2 briefly introduces reinforcement learning, explaining the important concepts and its main goals. It will also describe some of the recent advances inside of the field and the achievements of these methods.

Chapter 3 addresses Polyhedra analysis. It will explain its usefulness inside of static analysers and how it works. It will also introduce some of the modifications made to the domain, in order to reduce some of its shortcomings.

Chapter 4 addresses the work done throughout this thesis. It will first address how the two previous chapters were combined and the basic algorithm was designed. It will also show the different methods tested to further optimise the algorithm.


Chapter 5 discussed the different experiments that were run to find the optimal combination of the different parts of the algorithm. It will also present the results of the execution of the finalised algorithm on different benchmarks.

Chapter 6 will be a final discussion of the work achieved throughout this thesis.

Reinforcement Learning

2

Reinforcement Learning is a very general problem of machine learning studied in a multitude of different fields, such as game theory, information theory or statistics. It is a very broad concept and the foundations are the following. An agent interacts with a given environment, at its disposal, it has a set of various different actions. When an action is performed it receives a reward. The goal of the agent is to devise an action selection strategy that should maximise the cumulative reward of the actions.



2.1. Concepts

In order to be able to solve a problem with reinforcement learning, this problem has to be mapped to the following RL concepts:

- A set of agent states S , with the initial state $s_0 \in S$
- A set of action A
- A function giving the reward of performing an action from s_t to s_{t+1} , $r(s_t, a_t, s_{t+1}) \in \mathbb{R}$


During the execution of the program, the agent will first start in the initial state s_0 . Then, at each timestep $t = 0, 1, \dots$, the agent will pick an action $a_t \in A$ then the action will be executed the agent will move from s_t to s_{t+1} and receive the reward $r(s_t, a_t, s_{t+1})$. This process will be repeated until a final state is reached. In RL, state transitions typically satisfy the Markov property. This property states that the next state s_{t+1} only depends on the current state s_t and on the action taken at this state a_t . We call the sequence of actions and states from the initial to

2. Reinforcement Learning

the final state an episode. The goal of the agent is to devise an action picking strategy so as to maximise the cumulative reward. I will demonstrate the idea of the cumulative reward with the following example.

Example 2.1

Let's assume we have an episode with two time-steps and an action set with only two possible actions. In the following tree, we can see the different actions we can take at each state and their respective reward.



The cumulative reward, is the summed reward of all the actions taken during an episode. This is shown at the end of the leaf nodes in the example above. Therefore, an ideal decision policy would choose the actions $\{a_1, a_0\}$ in that order for the above episode.

2.2. Q-function

We call Q-function or quality function, a mapping $Q : S \times A \rightarrow \mathbb{R}$ that specifies the cumulative reward of picking an action a_t in state s_t . If this function was known the reinforcement learning problem would become quite trivial, as it would simply suffice to pick, at every time step, the action with the highest Q-value.

Unfortunately, as in most real world cases the state space is very large or even infinite. Computing the Q-function exactly is very close to infeasible.

2.3. Q-function approximation

The goal of reinforcement learning is to obtain the best possible Q-function approximation. To achieve this, the agent is allowed to interact freely inside of the environment, picking actions randomly or with the till now learned policy. It observes the different results it obtains and updates its policy accordingly.

I shall now talk about a few concepts used during training.

Exploration-exploitation

As mentioned above, during training we either pick the best action with the till-now learned policy or we just pick a random action. The trick is in finding a correct balance between the both of these as both of these have their advantages. If we always pick the best action we will converge quickly but the risk of getting stuck in a local minima is fairly high. On the other hand, if we only pick random actions, the convergence will be very slow and we might simply diverge and not be able to find a strategy at all. Most training algorithms start with a high probability of exploration and then as training progresses increase the exploitation probability as to solidify the learned policy.

Learning rate

The learning is also part of the exploration-exploitation dilemma. With this ratio we can modify the importance that we give to newly acquired information. The lower it is the harder it will be to overwrite a already learned policy.

Discount factor

In reinforcement learning, we approximate the Q-function with the following equation:

$$Q'(s, a) = \max_{\pi} \mathbb{E}[r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} | s_t = a, a_t = a, \pi] \quad (2.1)$$

where π is the learned policy.

We call γ the discount factor. It represents the importance that we give to the future reward. If $\gamma = 0$ then only the immediate reward will be considered and if $\gamma \approx 1$ then we will look for a policy that will maximise the cumulative reward. The same as before one has to choose between a faster convergence and a higher risk of getting stuck in local minima.

2.3.1. Q-learning

Q-learning is one of the main algorithms used in reinforcement learning in order to approximate the Q-function. It models the Q-function as a set of basis function where each basis function assigns a value to a (state,action) pair and each feature has its own basis function. Q-learning offers several advantages. It is efficient and converges relatively quickly. The learned policy can be quite easily interpreted. By definition, Q-learning uses a linear function approximation. In most cases this should not cause much of a problem. But, in the case where the ideal decision policy is non-linear, Q-learning will never achieve optimal results.

2.3.2. Deep Q-networks

Deep Q-networks are a novel method used for reinforcement learning that has gained a lot of attention in the recent years. The reason why it has gained so much notice lately, is due to the

2. Reinforcement Learning

ability of a single algorithm being able to achieve very good results on a broad array of tasks without the need for any specialized knowledge about the task beforehand, leading some to call it as the first major steps towards general artificial intelligence.

What separates deep Q-networks from other reinforcement learning methods, is that they use neural network in order to approximate the Q-function. Neural networks are non-linear function, reinforcement learning has been known to diverge when non-linear function approximators have been used. However, the development of new technique such as experience replay memory has allowed them to become a viable tool for reinforcement learning and to achieve some very remarkable results, most notably in domains such as video games.

Polyhedra Analysis

3

Polyhedra analysis is one of the main tools of static analysis. During the execution of a program variables can become bounded, either by numbers or by other variables. Polyhedrons are one of the most expressive ways of modelling all the possible values that variables can have as well as the different dependencies between them. Different alternatives exist for constraint representation instead of polyhedra, but the polyhedra domain is by far the most expressive. Unfortunately, it has worst case exponential space and time complexity meaning that using it on most real-world applications would cause it to either timeout or to run out of memory making it very impractical. Therefore, other domains have been more widely used such as octagon, zone or pentagon, but all of these are less expressive and therefore less precise by design. In the recent years several techniques have been developed that have managed to speed up polyhedra analysis without the loss of precision.

3.1. Polyhedra representation


One of these techniques involves separating the way we represent our polyhedra. Polyhedra can be represented with both their constraint representation and their generator representation. To illustrate these different illustrations I will proceed with the following example. Lets assume we have the following set of assertions:

Example 3.1

```
if  $x \geq 0.5 \wedge x \leq 1.5 \wedge y \geq 0.5 \wedge y \leq 1.5$  :  
    ...  
end
```

Inside of the if statement, the polyhedron would have the following shape:

3. Polyhedra Analysis



We can represent this information in two different possible ways.

3.1.1. Constraint representation

In constraint representation we model the polyhedron as the intersection of a finite number of closed half spaces and a finite number of subspace. The resulting polyhedron can be written as:

$$P = \{x \in Q^n | Ax \leq b \wedge Dx = e\} \quad (3.1)$$

where A,D are matrices and b,e are vectors of natural numbers. Therefore, the constraint representation of the above example would be:

$$C = \{-x \leq -0.5, x \leq 1.5, -y \leq -0.5, y \leq 1.5\} \quad (3.2)$$

3.1.2. Generator representation

In order to encode a Polyhedron with the generator representation, we have to model it as the convex hull of three items:

- Set of vertices $V \in Q^n$.
- Set off rays where one end is bounded and that start from a vertex, modelling the edges of the polyhedron.
- set off lines modelling the infinite edges of the polyhedron with both ends unbounded.

The result of generator representation on the previous example would have the following form:

$$G = \{V = \{(0.5, 0.5), (1.5, 0.5), (1.5, 1.5), (0.5, 1.5)\}, R = \emptyset, Z = \emptyset\} \quad (3.3)$$

3.2. Polyhedra domain

Now that we can represent our Polyhedra we can do some interesting calculations with them. The polyhedra abstract domain consists of the polyhedral lattice: $(P, \sqsubseteq, \sqcup, \sqcap, \perp, \top)$, and a set of operators that we can apply onto the different polyhedra. The different operators are the following:

- Inclusion test: $P \sqsubseteq Q$
- Equality test: $P = Q$
- Join: $P \sqcup Q$
- Meet: $P \sqcap Q$
- Widening, this operator is applied to accelerate convergence since the polyhedral lattice has infinite height:

$$C_{P \nabla Q} = \begin{cases} C_Q & \text{if } P = \perp; \\ C'_P \cup C'_Q, & \text{otherwise;} \end{cases}$$

where $C'_P = \{c \in C_P \mid C_Q \vdash c\}$, and

$C'_Q = \{c \in C_Q \mid \exists c' \in C_P, C_P \vdash c' \text{ and } ((C_P) \cup \{c\}) \vdash c'\}$ where $C \vdash c$, test whether c can be entailed from constraints in C

- Conditional: let $\otimes \in \{\leq, =\}$, $1 \leq i \leq n$, $\alpha \in Q$ then $\alpha x_i \otimes \delta$ adds the constraint $(\alpha - a_i)x_i \otimes \delta - a_i x_i$ to the constraint set C
- Assignment: $x_i = \delta$, first adds x_i to P then augments C with $x_i - \delta = 0$

In the following table we can see the respective complexities of the different operators according to the representation

Operator	Constraint	Generator	Both
Inclusion (\sqsubseteq)	$O(mLP(m, n))$	$O(gLP(g, n))$	$O(ngm)$
Join (\sqcup)	$O(nm^{2^{n+1}})$	$O(ng)$	$O(ng)$
Meet (\sqcap)	$O(nm)$	$O(ng^{2^{n+1}})$	$O(nm)$
Widening (∇)	$O(mLP(m, n))$	$O(gLP(g, n))$	$O(ngm)$
Conditional	$O(n)$	$O(ng^{2^{n+1}})$	$O(n)$
Assignment	$O(nm^2)$	$O(ng)$	$O(ng)$

$m = |C|, g = |G|, LP(m, n)$ is the complexity of solving a linear program with m constraints and n variables

As we can see no representation is faster than the other. As some operators are quicker in one but others in the other. But, as we can see in the last table, when both representations are available all operators are polynomial.

Chernikova's Algorithm

The first optimisation that one can do is keep both representation of the the polyhedron and

3. Polyhedra Analysis

for each operator picking the representation that minimises the time complexity. A conversion between the two representations is possible thanks to Chernikova's algorithm.

3.3. Polyhedra Decomposition

Another technique used to increase the efficiency of polyhedra analysis, is that of online decomposition. It is based on the observation that during the execution of a program, not all its variables are dependent on one another. Using this observation, we can separate the set of all variables into independent sets. Therefore, instead of having to represent the whole set of variables with one large polyhedron we can instead represent it with various smaller ones.

Let's assume we have a set of variables χ in a Polyhedron P . The set χ can be partitioned as $\pi_P = \{\chi_1, \chi_2, \dots, \chi_r\}$, $\chi_i \subseteq \chi$. We call the partitioning of the set permissible iff $\chi_i \cap \chi_j = \emptyset$, $\forall i \neq j$. Once the decomposition has been done in this way, during the execution of an operator, it only has to be executed on the subset of blocks that are influenced by it. This allows for a very large performance gain. Giving us the following time complexity for the various operators.

Table 2 Asymptotic time complexity of Polyhedra domain operators with decomposition

Operator	Decomposed
Inclusion (\sqsubseteq)	$O(\sum_{i=1}^r n_i g_i m_i)$
Join (\sqcup)	$O(\sum_{i=1}^r n_i g_i m_i + n_{max} g_{max})$
Meet (\sqcap)	$O(\sum_{i=1}^r n_i m_i)$
Widening (∇)	$O(\sum_{i=1}^r n_i g_i m_i)$
Conditional	$O(n_{max})$
Assignment	$O(n_{max} g_{max})$

Example 3.2

Let's consider the with variables $X = \{x_1, x_2, x_3, x_4\}$ and $C = \{x_1 + 2 \cdot x_3 \leq 10, x_2 = 1\}$. This polyhedron can be partitioned into three blocks $\pi_P = \{\{x_1, x_3\}, \{x_2\}, \{x_4\}\}$ and its corresponding factors are equal to:

$$C_{P_1}\{x_1 + 2 \cdot x_3 \leq 10\}, C_{P_2}\{x_2 = 1\}$$

3.3.1. Decomposition operators

As we change the model of our domains we must equally update the operators inside of these domains. For sake of brevity I will not go into detail about each of them, but I will only talk about the join operator as it plays an important role in the work of this paper.

During the join of P and Q , most of the time their factors will not be equal ($\pi_P \neq \pi_Q$). Therefore, we have to remake their partitions in order for their factors to be equal $\pi = \pi_P \sqcup \pi_Q$. During most joins of a normal execution this will not cause much problem, but during some cases the join can merge all blocks producing the \top partition. In order to rebuild the \top partition from all its blocks we use the following formula:

$$P = P_1 \bowtie P_2 \bowtie \dots \bowtie P_r = (C_{P_1} \cup C_{P_2} \cup \dots \cup C_{P_r}, G_{P_1} \times G_{P_2} \times \dots \times G_{P_r}) \quad (3.4)$$

Due to the cartesian product, building the \top partition can blow up the number of generators and therefore cause online decomposition to loose its performance advantage.

3.4. Reinforcement Learning for polyhedra analysis

Both of the techniques presented in chapters 3.2 and 3.3, manage to achieve considerable performance gain without having to sacrifice any precision. Unfortunately, at some point compromises have to be made. Many different algorithms have been proposed that try their best at minimising the precision loss and maximising the resulting performance gains.

As explained in 3.3.1. one bad constraint can significantly decrease the performance of the whole program. The trick is being able to identify this variable at the correct time. One of the potential solutions to this problem that we shall explore in this paper is training a reinforcement learning algorithm in order to decide when to apply abstractions.

3.4.1. Adapting polyhedra analysis for Reinforcement Learning

In order to be able to apply reinforcement learning, we first have to express polyhedra analysis in terms of reinforcement learning. That is, we need an agent, a set of actions, a reward and a set of features.

Conceptually the problem is quite simple. We will have our agent, the static analyser, executing the analysis of a program. During analysis it will have a set of actions at its disposal. Each of these actions will have a varying performance/precision capability. The goal of the agent will be to choose the correct degree of abstraction at the correct point in time. The reward has to be some sort of compromise between the precision and the performance of the given action.

3.4.2. Existing methods

Existing methods already exploit this idea. Using techniques such as Q-learning in order to create a decision policy. Q-learning is a common linear Q-function approximation technique. However recent RL methods concentrate on the use of non-linear Q-function approximation such as neural networks. Whilst non-linear Q-function approximations have been known to be unstable or divergent in the past, development of new techniques such as experience replay has helped make them a viable tool for a series of reinforcement learning problems. In the following chapters, we shall explore the idea of using such a technique for polyhedra analysis.

3. *Polyhedra Analysis*

Fast polyhedra analysis with deep Q-networks

4

The goal of my work was to build upon the existing solutions, used to make polyhedra analysis more efficient and use deep Q-networks to further improve the performance. Specifically, my goal was to extend the Elina Framework, the version of it that used reinforcement learning for polyhedra analysis, in order to use deep Q networks for the Q function estimation. The work I have done can be separated into three major subtasks. First of all, I had to figure out a way of calling a deep neural network from Elina, so that I could use it for the Q function estimation. Secondly, In order to test out the validity of deep Q-networks on polyhedra analysis and to have a baseline for further optimisation. I implemented the basic deep Q-network training algorithm. Finally, I attempted to further optimise this problem with some task specific knowledge.

4.1. Incorporating Neural Networks inside Elina

The first part, whilst definitely being the least interesting was probably the most laborious and a very key part of the whole work. The problem was that for my Q value estimation I wanted to use neural network regression from Keras in Python. I decided to use this framework as it is very widely used for this type of problem, is relatively easy and intuitive to use whilst remaining very powerful. However, Elina being written in C, I somehow had to make the link between my Python code and the C code of Elina. I finally achieved this by using the Python/C API and embedding the python code into the Elina framework. Initialising all the python code as well as importing all the necessary libraries such as Keras, is done in the `op_pk_manager_alloc()` function so that it doesn't have to be done every time we call the join function. Inside the join function we simply have to call the learning or prediction algorithms.

4.2. Basic algorithm

In order to get the basic training algorithm working, first the domain of polyhedra analysis has to be made compatible with the domain of reinforcement learning. Therefore the following list of items have to be initialised inside of the polyhedra domain:

- A set of features describing the polyhedra to use as states S
- A set of actions A

4. Fast polyhedra analysis with deep Q-networks

- A reward function r

Luckily, since reinforcement learning methods have already been used for polyhedra analysis we can simply reuse what has already been developed for our baseline algorithm.

4.2.1. Features

In the Reinforcement learning method the following nine features was used.

Feature	Extraction complexity	Approximate range	Scaling
$ \beta $	$O(1)$	1-10	$x/1.$
$\min(\chi_k : \chi_k \in \beta)$	$O(\beta)$	1-50	$\text{round}((x/5), 0.5)$
$\max(\chi_k : \chi_k \in \beta)$	$O(\beta)$	1-50	$\text{round}((x/5), 0.5)$
$\text{avg}(\chi_k : \chi_k \in \beta)$	$O(\beta)$	1-50	$\text{round}((x/5), 0.5)$
$\min(\bigcup G_{P_m}(\chi_k) : \chi_k \in \beta)$	$O(\beta)$	1-10000	$\text{round}((x/1000), 0.1)$
$\max(\bigcup G_{P_m}(\chi_k) : \chi_k \in \beta)$	$O(\beta)$	1-10000	$\text{round}((x/1000), 0.1)$
$\text{avg}(\bigcup G_{P_m}(\chi_k) : \chi_k \in \beta)$	$O(\beta)$	1-10000	$\text{round}((x/1000), 0.1)$
$\{ x_i \in X : x_i \in [l_m, \infty) \text{ in } P_m \} +$ $\{ x_i \in X : x_i \in (-\infty, u_m] \text{ in } P_m \}$	$O(n_g)$	1-50	$\text{round}((x/5), 0.5)$
$\{ x_i \in X : x_i \in [l_m, u_m] \text{ in } P_m \}$	$O(n_g)$	1-50	$\text{round}((x/5), 0.5)$

4.3. Actions

As we discussed in 3.3.1., the bottleneck of the decomposed analysis is the join. Therefore, the set of actions will be a various set of joins of different precision and performance.

First of all, the cost of the joins depends on the size of the block. Therefore, bounding the the size of the block with a threshold would increase the performance. These four different thresholds are used $\text{threshold} \in [5, 9], [10, 14], [15, 19], [20, \infty)$

Once we have decided on the size of a threshold we have to equally decide on what to do if the block has a greater size than the threshold. Their are different possibilities of how to split a large block into smaller ones, but basically, one has to pick a subset of constraints to remove from the block so that all its constraints are no longer dependant and then it can be further partitioned. The following three constraint removals are used:

Stoer-Wagner min-cut

This uses the basic idea of removing the minimal amount of constraints in order to be able to

split to block into two separate permissible partitions.

Weighted constraint removal

The second and third constraint removal techniques are based on the same principal. They associate a certain weight to each constraint and then remove the constraint with the highest weight. Two different weight distribution techniques are considered.

In the first one, we first compute for each variable $x_i \in X$ the number of constraints it appears in, we can call this n_i . Then for each constraint c_i we set its weight to the sum of n_i of all the variables that are in the constraint.

The second method, we first compute for each pair of variables $x_i, x_j \in X$, n_{ij} the number of constraints containing both x_i and x_j . The weight of the constraint is then the sum of all the n_{ij} of all the pairs of constraints x_i, x_j contained in the constraint.

The idea of removing the constraint with maximal weight, is that, most likely the variables occurring in the constraint are also bounded by other constraints and therefore will not become unbounded once the constraint is removed.

Merging blocks

The next three actions are various block merging strategies. The idea is to select different blocks and merge them together as long as the resulting blocks remain below the threshold in order to increase the precision of the subsequent join. The following three block merging strategies are used:

- No merge, no blocks are merged
- Merge smallest first, we first merge the smallest two blocks together. We then remove the smallest block and continue as long as the resulting merge remains below the threshold.
- Merge small with large, similar to the previous strategy but this time we merge the smallest block with the largest.

In total we have four different thresholds, three different constraint removal algorithms and three different block merging strategies. We can mix and match these together as we please, which means that in total we have $4 \times 3 \times 3 = 36$ different actions we can pick.

Once this has been done, I could finally proceed to the more interesting part of writing the training algorithm. Under its most basic form, the training algorithm could simply be the following. Pick a random reward or maximal action, observe the reward and then retrain the network with the observed plus the discounted future reward on the given action. This very basic algorithm however did not get me very far as I ran into two main problems.

4.4. Reward

As a reminder, the objective of the reward is to guide our learning policy. Rewarding it when it takes actions towards our global goal and penalising it when it does otherwise. Therefore, the

4. Fast polyhedra analysis with deep Q-networks

reward developed by the Q-learning method was the following:

$$r(s_t, a_t, s_{t+1}) = 3 \cdot n_s + 2 \cdot n_b + n_{hb} - \log_{10}(cyc) \quad (4.1)$$

Where n_s is the number of exactly defined variables of the resulting polyhedron after the join. n_b is the number of bounded variables and n_{hb} the number of semi-bounded variables of the resulting polyhedron. cyc is the number of cycles required to perform the join.

Back to deep Q-networks

As we have finally initialised polyhedra analysis for the use with reinforcement learning. We can finally proceed to the more interesting part of designing the training algorithm. Under its most basic form, the training algorithm could simply be the following. Pick a random reward or maximal action, observe the reward and then retrain the network with the observed plus the discounted future reward on the given action. This very basic algorithm however did not get me very far as I ran into two main problems.

Divergence of decision policy

The first of which was the divergence of the action selection strategy. Meaning that neural networks were not able of creating a consistent strategy but would be picking different actions all the time.

Divergence of the action value prediction

The second problem was that the predictions of my estimator would diverge towards infinity. This is also known as the exploding gradient problem.

These problems are not new and are a fundamental problem of trying to use non-linear functions for the Q-function approximation. Several techniques exist to combat these problem. Many of which reduce the complexity of the network and or the problem in general. However, some new methods have been developed recently that would allow the training of the estimators whilst retaining the problems original complexity. I decided the use the following two concepts inside of my algorithm.

4.4.1. Experience replay memory

Experience replay memory is a biologically inspired mechanism. That gives the estimator a very basic concept of a memory and instead of directly learning from the current events happening. The agent learns from a random subset of its memroy. More formally, during training, an array of a certain size filled with the past memory objects is kept. Each memory item contains the following items: $mem(s_t, a_t, r_t, s_{t+1})$. The current state, the action taken at this state, the observed reward and the next state. For training we then simply pick a random subsample from the memory array and train from these examples.

The objective of the memory is to homogenise the training data. It comes from the observation that during the execution of a program, often a particular strategy would be optimal during a certain period of time and afterwards another one would be the new optimal. This caused two

major problems. First of all, it is not very time efficient as we do not gain much information by learning from the same data. Secondly, as there was low diversity in the training data and the decision strategy would frequently abruptly change it. This either led to the neural network getting stuck in local minima or simply to diverge and not obtain a policy.

Picking a random subsample from memory helps increase the variance amongst the training data allowing the network to learn a more global policy.

4.4.2. Separating target from max Q estimators

The other method used to reduce the divergence of the predicted Q values towards infinity, was to reduce the correlation between the training and the prediction data. This was done because, since the networks were being fitted based on the maximum Q value estimation, diverging towards infinity reduced their error and therefore was a valid strategy that they would use. In order to reduce this correlation, the networks predicting the maximum Q value and the ones predicting the Q value were separated. The weights of the maximum Q-value estimators were then updated every n steps in order for them to remain up to date.

Further modifications were also made on the neural network level in order to reduce these problems. I will get into these in part.

Once these modifications have been made the basic training algorithm has the following form.

Pseudo code of basic learning algorithm

4. Fast polyhedra analysis with deep Q-networks

Algorithm 1: DQN Training algorithm

1 **function** learn ($S, A, r, \gamma, \alpha, \phi, N, l_freq, b_size, update_nn_freq, \varepsilon$);

Input :

$S \leftarrow states, A \leftarrow Actions, r \leftarrow reward,$
 $\gamma \leftarrow discount\ factor, \alpha \leftarrow learning\ rate,$
 $\phi \leftarrow$ set of feature functions over S and A
 $N \leftarrow$ size of memory, $l_freq \leftarrow$ learning frequency,
 $b_size \leftarrow$ batch size, $\varepsilon \leftarrow$ random action probability,
 $update_nn_freq \leftarrow$ frequency of updating max Q estimators

Output:

$\theta \leftarrow$ trained weights of the estimator

2 $\theta =$ initialise random weights and learning rate α

3 $\theta_{max} =$ initialise random weights and learning rate α

4 $M =$ initialise an empty memory with size N

5 **for** each episode **do**

6 start from initial state $s_0 \in S$

7 **for** $t = 0, 1, 2, \dots$ **do**

8 Initialise new memory item m_t

9 With probability ε take random or maximal action a_t

10 observe next state s_{t+1} and $r_t(s_t, a_t, s_{t+1})$

11 Set $m_t.a = a_t, m_t.s_1 = s_t, m_t.s_2 = s_{t+1}, m_t.r = r_t$

12 **if** $M_{size} \geq N$ **then**

13 del M_0

14 $M_N = m_t$

15 **else**

16 push m_t on M

17 **if** $t \bmod l_freq = 0$ **then**

18 select a random batch of size b_size from m

19 compute $Q(:, s_t)$ estimation with θ

20 compute $Q(:, s_{t+1})$ estimation with θ_{max}

21 set $Q(a, s_t) = Q(a, s_t) + \gamma * max(Q(:, s_{t+1}))$

22 Fit weights θ with new training data

23 **if** $t \bmod update_nn_freq = 0$ **then**

24 set $\theta_{max} = \theta$

25 **end**

26 **end**

27 **return** θ

4.5. Further optimisations

The results of this algorithm will be discussed in the results section. I will now go into more details about further optimisations and different strategies that I tried in order to further improve the efficiency of the algorithm. I will now go into more detail about these modifications.

4.5.1. Separating the problem into two

The first problem specific modification I made, was subdividing the problem into two independent subproblems. The objective of Polyhedra Analysis is to obtain the most precise result in the quickest way possible. These are two independent objectives. My goal was therefore to separate these two tasks and create two independent subsystems, one focusing on making the results as precise as possible, and the other on the time complexity of getting there.

Such a division of the problem, offer two main advantages. Firstly, a greater flexibility during training. The parameters of each system can be tuned for its specific needs. So for example the characteristics of the neural network or the discount factor. Furthermore, the optimal features used for precision and performance estimation, are quite probably different. The separation therefore allows us to use only the necessary features for each subsystem making the learning and prediction less complex and therefore more precise and converge faster. The reward as well can be fine tuned for the needs of the specific needs of the subsystem.

Secondly, two separate subsystems also allow for a more flexible algorithm post training. Since during training, we will have trained two different Q function estimators, during prediction we will have a greater possibility of optimising our results by changing the importance we give to each subsystem at different points in time according to the specific needs. I will get into more detail about this in the section about the action selection policy.

The pseudo code for the training algorithm with separated estimators has the following form:

Pseudo code of the separated training algorithm

4. Fast polyhedra analysis with deep Q-networks

Algorithm 2: DQN Training algorithm

```

1 function learn ( $S, A, r, \gamma, \alpha, \phi, len, l\_freq, b\_size, update\_nn\_freq$ );
   Input :
        $S \leftarrow states, A \leftarrow Actions, r \leftarrow reward,$ 
        $\gamma \leftarrow discount\ factor, \alpha \leftarrow learning\ rate,$ 
        $\phi \leftarrow$  set of feature functions over  $S$  and  $A$ 
        $len \leftarrow$  size of memory,  $l\_freq \leftarrow$  learning frequency,
        $b\_size \leftarrow$  batch size,
        $update\_nn\_freq \leftarrow$  frequency of updating max Q estimators

   Output:
        $\theta_1 \leftarrow$  trained weights of neural network for performance
        $\theta_2 \leftarrow$  trained weights of neural network for precision

2   $\theta_1 =$  initialise random weights
3   $\theta_2 =$  initialise random weights
4   $\theta_{1\_max} =$  initialise random weights
5   $\theta_{2\_max} =$  initialise random weights
6   $m =$  initialise an empty memory
7  for each episode do
8      start with initial states  $s_{pr\_0} \in S, s_{pe\_0} \in S$ 
9      for  $t = 0, 1, 2, \dots$  do
10         Initialise new memory item  $m_t$ 
11         Take action  $a_t$  according to the action selection algorithm
12         observe next state  $s_{pr\_t+1}, s_{pe\_t+1}$  and  $r_{pe}(s_{pe\_t}, a_t, s_{pe\_t+1}), r_{pr}(s_{pr\_t}, a_t, s_{pr\_t+1})$ 
13         Set  $m_t.a = a_t, m_t.s_{pr\_1} = s_{pr\_t}, m_t.s_{pe\_1} = s_{pe\_t}, m_t.s_{pr\_2} = s_{pr\_t+1}, m_t.s_{pe\_2} =$ 
            $s_{pe\_t+1}, m_t.r_{pr} = r_{pr}, m_t.r_{pe} = r_{pe}$ 
14         if  $m_{size} \geq len$  then
15             del  $m_0$ 
16              $m_{len} = m_t$ 
17         else
18             push  $m_t$  on  $m$ 
19         if  $t \bmod l\_freq = 0$  then
20             select a random batch of size  $b\_size$  from  $m$ 
21             compute  $Q(:, s_t)$  estimation with  $\theta_1, \theta_2$ 
22             compute  $Q(:, s_{t+1})$  estimation with  $\theta_{1\_max}, \theta_{2\_max}$ 
23             set  $Q(a, s_t) = Q(a, s_t) + \gamma * max(Q(:, s_{t+1}))$ 
24             Fit weights  $\theta_1, \theta_2$  with new computations from this batch
25         if  $t \bmod update\_nn\_freq = 0$  then
26             set  $\theta_{1\_max} = \theta_1, \theta_{2\_max} = \theta_2$ 
27         end
28     end
29 return  $\theta_1, \theta_2$ 

```

4.5.2. Feature selection

Before I get into the specific features I used for my estimators, I would like to go a bit more in-depth about the theory of features and states. States in the case of polyhedra analysis would be the concrete Polyhedra. However, during our analysis we do not use the concrete polyhedra as this would be too complicated, but rather some information that we extract from them, we call this the feature vector. As explained in section 2.1., for reinforcement learning, states should respect the Markov property that states that s_{t+1} only depends on s_t and on a_t . Whilst this is the case for the polyhedra, our feature vectors do not have to obey this property as they do not represent the polyhedra themselves but only some information about them. The more precisely we manage to describe our polyhedron the closer we will be to respecting the Markov property and the better will be the results of RL.

With regards to the old set of features. One big inconvenience of using Q-learning, is that, since we want to represent our Q-function with basis functions, the size and dimensions of our feature vector is very limited. This constraint no longer holds when using more advanced methods such as deep Q-networks. This allowed me to make two major changes about the feature vector.

Adding new features

The first change I was able to make, was to add some new features. First of all, I reused the nine features from the Q-learning algorithm. The first seven features of these features are used to characterise the complexity of the join. They are the number of blocks, minimal, maximal and average size of the blocks and the minimal, maximal and average size of the generator set. The last two features are used to characterise the precision of the inputs and they are the number of variables with a finite upper and lower bound, as well as the number of variables with a finite upper or lower bound, in both Polyhedra.

I further extended this set with four new features, in order to further increase the description accuracy of the feature vector. The selection of these features was done by trial and error, with an experimental observation of an increase, or lack thereof, of accuracy. It is also possible to check whether a particular feature is useful once training is finished by analysing the neural network and observing the impact a particular feature has on the end result. However, since at the end my network had a total of four layers, this made its analysis somewhat complicated.

At the end I added a total of four new features, three of which are used for modelling the precision and one for the complexity. For the complexity I added the number of constraints. As for the precision I added the number of unconstrained variables, the number of exactly constrained variables and finally, the sum of the values the bounded variables can have (i.e. an approximation of the circumference of the polyhedra). I also tried using some features that would have perhaps modelled the precision more accurately, such as most notably approximating the volume of the polyhedra. Unfortunately, one must also consider the complexity of computing the features. The computation of the volume approximation was far too time complex, which greatly increased the learning time making the feature not viable.

Bucketing

Due to the limitations of Q-learning the old algorithm used a very restrictive bucketing policy.

4. Fast polyhedra analysis with deep Q-networks

Deep reinforcement learning does not suffer from these restrictions. I was therefore able to remove the bucketing from the feature selecting. I still implemented a version of bucketing in my algorithm for three main reasons.

Firstly, for features with very big values an exact precision is not needed. For example, for the feature that approximates the circumference of the polyhedra, its total value can be very big and if it is a million and one or just a million doesn't impact the resulting precision much therefore bucketing makes sense.

The second reason, for the use of bucketing is that it greatly decreases the learning time and helps convergence.

Finally, the last reason is that the decision policy doesn't give more importance to some features simply because they are larger than others. In my total of thirteen different features, they all can have very different values. For example, the number of blocks varies mainly between one and ten, and the circumference of the polyhedra can go up to 10^9 . This does not mean that the circumference has a greater impact on the overall precision of the polyhedra. Bucketing assures that all features have a similar impact on the decision policy.

In order to decide the values of the bucketing I would use, I ran my algorithm on a big number of benchmarks computing the maximal, minimal and average values the different features could have. Finally, I scaled them so that they would all approximately remain between the bounds of zero and ten. I then bucketed them in the following fashion:

Feature	Extraction complexity	Approximate range	Scaling
$ \beta $	$O(1)$	1-10	$x/1.$
$\min(\chi_k : \chi_k \in \beta)$	$O(\beta)$	1-50	$\text{round}((x/5), 0.5)$
$\max(\chi_k : \chi_k \in \beta)$	$O(\beta)$	1-50	$\text{round}((x/5), 0.5)$
$\text{avg}(\chi_k : \chi_k \in \beta)$	$O(\beta)$	1-50	$\text{round}((x/5), 0.5)$
$\min(\bigcup G_{P_m}(\chi_k) : \chi_k \in \beta)$	$O(\beta)$	1-10000	$\text{round}((x/1000), 0.1)$
$\max(\bigcup G_{P_m}(\chi_k) : \chi_k \in \beta)$	$O(\beta)$	1-10000	$\text{round}((x/1000), 0.1)$
$\text{avg}(\bigcup G_{P_m}(\chi_k) : \chi_k \in \beta)$	$O(\beta)$	1-10000	$\text{round}((x/1000), 0.1)$
$\{ x_i \in X : x_i \in (-\infty, \infty) \text{ in } P_m \}$	$O(n_g)$	1-100	$\text{round}((x/10), 0.5)$
$\{ x_i \in X : x_i \in [l_m, \infty) \text{ in } P_m \} +$ $\{ x_i \in X : x_i \in (-\infty, u_m] \text{ in } P_m \}$	$O(n_g)$	1-50	$\text{round}((x/5), 0.5)$
$\{ x_i \in X : x_i \in [l_m, u_m] \text{ in } P_m \}$	$O(n_g)$	1-50	$\text{round}((x/5), 0.5)$
$\{ x_i \in X : x_i \in [u_m, u_m] \text{ in } P_m \}$	$O(1)$	1-200	$\text{round}((x/20), 0.2)$
$ X $	$O(n_g)$	1-50	$\text{round}((x/5), 0.5)$
$\sum(u_m - l_m) : u_m, l_m \in \{[l_m, u_m] \text{ in } P_m\}$	$O(n_g)$	$2 * 10^9$	$\text{round}((x/2 * 10^8), 0.01)$

One thing to note is, that as opposed to the reinforcement learning algorithm, my features do not have a maximal possible value. I believe that this increases the precision of the q function

estimation, especially for the approximation of the complexity. Certain features can have very large values and I believe when this arrives, it has a major impact on the complexity of the join. However, these very large values were ignored by the bucketing of the reinforcement learning algorithm.

4.5.3. Reward Function modelling

Due to the nature of the new algorithm, the modelling of the reward function was separated into two separate subproblems.

Performance reward

Firstly, the reward for the precision estimator and then the reward for the performance estimator. Modelling the complexity is fairly simple and very straight forward, as simply counting the number of cycles needed for the computer to execute the join perfectly models the complexity and is exactly what we want to optimise. One thing to note is that we want this reward to be as small as possible, two simple solutions for this are either to invert it or to negate it. I experimented with both of these and found that negating it produces better results. I assume this is due to the fact that inverting the reward, makes it have a nonlinear curve and its derivative loses importance the higher the CPU cycles are, which is not something we want. Another thing to note is that modelling the complexity reward in this way gives us another advantage over the reinforcement learning version of the algorithm. The reinforcement learning algorithm took the \log_{10} off the CPU Cycles in order for the complexity and the precision reward to have similar values. I believe that this produces a similar problem as inverting the reward. The reward is no longer linear and its differences lose importance the higher it gets. Once again, this is not something that we want to model. The final reward has the following form:

$$r_{pr1}(s_t, a_t, s_{t+1}) = -1 \cdot cyc \quad (4.2)$$

Where *cyc* is the number of CPU cycles.

Precision reward

As to the second reward, modelling the precision of the resulting join. This is considerably more difficult than modelling the complexity as there is no trivial element giving us the precision of our polyhedron. In order to choose the reward, I proceeded by intuitively picking a small set of options and verifying them experimentally at the end. At the end, I tested out three different reward.

I took the first one from the reinforcement learning algorithm, that is, the objective is to maximise the amount of exactly bounded, bounded and half bounded variables. The second is a further extension by penalising the amount of values a bounded variable can have in the resulting Polyhedra. The final, reward function is slightly different. In this one, I penalise the loss of an exactly bounded, bounded or half bounded variable. Reward the loss of a bounded variable and penalise the amount of values a bounded variable can have. I also tried basing a reward function on an approximation of the volume of the resulting Polyhedra, unfortunately, the complexity of this computation was too high which made it too impractical to use in practice. The

4. Fast polyhedra analysis with deep Q-networks

final rewards have the following form:

$$r_{pr1}(s_t, a_t, s_{t+1}) = 3 \cdot n_s + 2 \cdot n_b + n_{hb} \quad (4.3)$$

$$r_{pr2}(s_t, a_t, s_{t+1}) = 3 \cdot n_s + 2 \cdot n_b + n_{hb} - \log_{10}(|n_b|) \quad (4.4)$$

$$r_{pr3}(s_t, a_t, s_{t+1}) = 3 \cdot (n_{sf} - n_{si}) + 2 \cdot (n_{bf} - n_{bi}) + (n_{hb_f} - n_{hb_i}) - \log_{10}(|n_b|) \quad (4.5)$$

Where n_s is the number of exactly defined variables of the resulting polyhedron after the join. n_b is the number of bounded variables and n_{hb} the number of semi-bounded variables of the resulting polyhedron. n_{sf} is the number of variables after the join and n_{si} before the join.

4.5.4. Action selection algorithms

As discussed before, I separated the Q function estimation into two separate subproblems. As I already mentioned earlier, this allows a certain amount of benefits that would not be possible otherwise. However, it also imposes one major complication. These two problems are not totally separable, as once we have predicted the two q function estimations, we have to somehow merge the information contained in both of these estimations and pick an ideal action accordingly. I tried out a few different action selection algorithms in order to find the one that maximises precision and performance the most. One thing to note is that it is at least partially possible to test out these algorithms post training. That is to say that we do not have to train using these in order to measure their performance afterwards. This only works if we train purely randomly however. As if, if we were to train using one selection algorithm and then test with another, this would surely deteriorate the results. The fact that we are able to test the selection algorithms after random training is still very helpful as the training time is relatively high and having to train for each algorithm would be very time intensive.

Algorithm 1

The first selection algorithm I used is perhaps the most intuitive one. First scaling both q function estimations. The performance one between $[-1,0]$ and the precision one between $[0,1]$. Adding the values together and picking the action with the maximal value. Whilst seeming very fair this type of selection has a couple of fundamental flaws. First of all, reinforcement learning estimates the best action to take in order to maximise the long-term objective, it however has less of a guarantee about the second to best and third to best action. This means that if the network is well trained, the chances that the action with the maximal Q-value prediction is also the best action to take at this point in time should be quite high. However, the guarantee that the action with the second highest Q-value prediction is the second-best action to take is much smaller. Using this selection algorithm tends to quite often not choose the best action but the second best or the third etc... making the probability that they are also good actions also a lot smaller. The second problem with this type of selection algorithm is that scaling the reward function causes a loss of information that could be very important. I will demonstrate this with an example, let's say we have four joins j1,j2,j3,j4. J1 takes one second, j2 ten seconds, j3 one hour and j4 ten hours. Let's also say that j1 and more precise than j2 and j3 is more precise than j5. This algorithm is going to treat the difference between j1 and j2 the same as the one between

j3 and j4, even though the gain in precision might be worth to loss in performance for the case of j1 and j2, whilst this would probably not be case for j3 and j4.

Pseudo code of algorithm 1

Algorithm 3: Action selection algorithm 1

```

1 function select1 ( $Q_{pr}, Q_{pe}$ );
   Input :
        $Q_{pr} \leftarrow$  array of Q-function estimations for precision,
        $Q_{pe} \leftarrow$  array of Q-function estimations for performance
   Output:
        $a \leftarrow$  action to take
2    $Q_{pr} = Q_{pr} / \max(Q_{pr})$ 
3    $Q_{pe} = Q_{pe} / \min(Q_{pe})$ 
4    $a = \max_{arg}(Q_{pr} + Q_{pe})$ 
5   return  $a$ 

```

Algorithm 2

The second algorithm used, that directly confront both problems of the previous algorithm proceeds as follows. We set some sort of threshold for the performance. We then look at the action that has the maximal Q-value for precision. If this actions Q-value for performance is above the threshold we take it otherwise we take the second-best action and continue until we find an action that is above the threshold. Again, this algorithm seems quite good at first glance and maybe if it was executed perfectly it would be the best option, but like the one before it has some fundamental problems. First of all, picking a threshold is not a very simple task. The q function estimator is a regression neural network, the activation function of its last layer is linear, this means that the output values are unbounded and they do not have a direct correlation with Realtime CPU cycles. In order to pick a threshold, I had to experimentally try different possible values and observe the resulting precision and adjust accordingly, however this threshold would vary according to the benchmark and therefore choosing an optimal one was a very difficult task on its own. Another problem of this algorithm is that If no action has a q performance estimation above the threshold the algorithm will choose the action with the worst precision, and this one does not even have to have the best performance estimation.

Pseudo code of algorithm 2

4. Fast polyhedra analysis with deep Q-networks

Algorithm 4: Action selection algorithm 2

```

1 function select2 ( $Q_{pr}, Q_{pe}, PE_{thresh}$ );
   Input :
        $Q_{pr} \leftarrow$  array of Q-function estimations for precision,
        $Q_{pe} \leftarrow$  array of Q-function estimations for performance,
        $PE_{thresh} \leftarrow$  performance threshold
   Output:
        $a \leftarrow$  action to take
2 while  $max(Q_{pr}) \neq 0$  do
3      $a = max_{arg}(Q_{pr})$ 
4     if  $Q_{pe}(a) \geq PE_{thresh}$  then
5         break;
6      $Q_{pr}(a) = 0$ 
7 end
8 return  $a$ 

```

Algorithm 3

Another possibility is to slightly modify the previous algorithm in order to minimise some of its short comings. This time we set both some sort of a threshold for the performance and a certain number x . if the action that has the highest precision estimation is under the threshold for performance, we take the x best actions according to their precision and take the one that has the highest performance estimation. The problem of picking an appropriate threshold remains the same and this time we have the further problem of picking a correct value for x . However, the case that all actions are under the threshold is no longer a problem. It is worth noting that this algorithm also has a greater time complexity, however this is still greatly outweighed if the correct action is taken.

Pseudo code algorithm 3

Algorithm 5: Action selection algorithm 3

```

1 function select3 ( $Q_{pr}, Q_{pe}, PE_{thresh}, N_{act}$ );
   Input :
        $Q_{pr} \leftarrow$  array of Q-function estimations for precision,
        $Q_{pe} \leftarrow$  array of Q-function estimations for performance,
        $PE_{thresh} \leftarrow$  performance threshold,
        $N_{act} \leftarrow$  number of actions to consider if below threshold
   Output:
        $a \leftarrow$  action to take
2    $a = \max_{arg}(Q_{pr})$ 
3   if  $Q_{pe}(a) \leq PE_{thresh}$  then
4        $a_{max} = a$ 
5        $val_{max} = Q_{pr}(a)$ 
6        $Q_{pr}(a) = 0$ 
7       for  $i \in N_{act}$  do
8            $a = \max_{arg}(Q_{pr})$ 
9           if  $Q_{pe}(a) \geq val_{max}$  then
10               $a_{max} = a$ 
11               $val_{max} = Q_{pe}(a)$ 
12               $Q_{pr}(a) = 0$ 
13       end
14   return  $a_{max}$ 

```

Algorithm 4

Finally, the last selection algorithm that I will talk about here is based upon the last one and made in such a way as to reduce the importance of correct parameter choosing. This time we begin by scaling the performance prediction to $[-1,0]$. We set a threshold to $[0,1]$. We pick the action that has the highest precision q estimation. If the absolute value of this action's performance estimation minus the maximal performance estimation is below the threshold, then we pick this action otherwise we pick the second-best precision action until we find one that is below the threshold. This time, since we compare to the best performance action we are guaranteed to be below the threshold at some point, in the worst case we will pick the action with the best performance. The threshold is also a lot easier to set since it is bounded. Unfortunately, once again we have the problem caused by the scaling of the performance estimation. However, after some experimental observation, I saw that when the join is fast all the q performance estimations tend to be quite close together, so hopefully this will not be all too much of a problem.

Pseudo code algorithm 4

4. Fast polyhedra analysis with deep Q-networks

Algorithm 6: Action selection algorithm 4

```
1 function select4 ( $Q_{pr}, Q_{pe}, PE_{thresh}$ );  
   Input :  
        $Q_{pr} \leftarrow$  array of Q-function estimations for precision,  
        $Q_{pe} \leftarrow$  array of Q-function estimations for performance,  
        $PE_{thresh} \leftarrow$  performance threshold  
   Output:  
        $a \leftarrow$  action to take  
2    $Q_{pr} = Q_{pr} / \max(Q_{pr})$   
3    $Q_{pe} = Q_{pe} / \min(Q_{pe})$   
4   while True do  
5        $a = \max_{arg}(Q_{pr})$   
6       if  $Q_{pe}(a) + PE_{thresh} \geq \max(Q_{pe})$  then  
7         break;  
8        $Q_{pr}(a) = 0$   
9   end  
10  return  $a$ 
```

4.6. Neural network characteristics

As for the characteristics of the neural network that I used. I started with a relatively small neural network and the grew it progressively as I expanded the number of vectors and the size that these can have. In the final version I use a fully connected neural network with four hidden layers of two hundred nodes each. They have each thirteen inputs, one for each feature and thirty-six outputs, one for each action. The first three layers use the Relu activation function and the last one uses a linear activation, since the rewards can be either negative or positive. I use stochastic gradient descent for updating the weights of the nodes and I clip its norm to 1. I do this in order to avoid the exploding gradients problem and prevent the predicted values to diverge towards infinity. I use the mean squared error in order to calculate the loss. I based the characteristics of my neural networks based on reading as these types of networks seem to be the most widely used ones for deep reinforcement learning. It is worth noting that I have not done much parameter optimisation on my neural networks as the training time of the algorithm is quite long which would make the parameter optimisation a very tedious task. It is also worth noting that I use the same network models for both performance and precision prediction.

4.7. Training

The training of the algorithm was done in multiple stages. During the first stage, only random actions were picked. This was done in order for the algorithm not to get stuck in local maxima but have the most global policy possible. I also did not have any time constraints about the length of training so I could afford to proceed in this way in order to get the best possible results. Throughout the next stages, the training was separated into two. One concentrating on

the precision and the other on the performance. All the actions were no longer chosen randomly, but now the action with the highest predicted value was chosen. The probability with which we chose a random action was progressively decreased throughout training. This was done in order to reinforce the choice of the best action and increase convergence speed. Once these stages finished, we would have two separate systems each optimised for their own task. In the final stage, I would then combine both of the systems and do a final training phase with the chosen action selecting algorithm as described above. I did this in order to further optimise the decision policy for the action selection algorithm.

There were two major complications that came up during training that had to be dealt with. The first of these was that, especially during random training the likelihood of getting stuck on some join and not finishing the analysis of a benchmark was quite high. To confront this I set a timeout of thirty minutes on each benchmark. The goal of this is to find a compromise between having enough time to get to interesting benchmarks and not wasting too much time in case the analysis gets stuck on a certain join. The other problem was that the analysis is very sensitive. That is to say, that one bad action can have severe consequences on the end precision of the analysis. This means that training with a low probability of picking random actions was very important for the estimators and had to be given enough time.

This method of training is quite time intensive. I proceeded in this way since time was not a big factor but the end results were more important. A more efficient training strategy can surely be found if this is needed.

4. *Fast polyhedra analysis with deep Q-networks*

Results

5

As described in the previous chapter, different strategies were elaborated for the reward function and the action selection algorithm. Training all possible combinations and then observing the end precision would be a viable way of finding the optimal combination of these. However, due to the training time of the algorithm, this would be very time intensive. I used the following techniques in order to find the best algorithms.

5.1. Reward function selection

As a reminder, if we view the problem as a decision tree, and the role of reinforcement learning to pick the best path at each node. The reward function can be viewed as the compass of the RL algorithm that should leading the algorithm to the end in such a way as to maximise the global goal.

I decided to measure the effectiveness of each reward function by observing if the direction it steered the algorithm to, was indeed correctly maximising this global goal. To achieve this, I would first train an algorithm with each of the different reward functions. Once this was done, I modified the Q-learning version of the algorithm so that during testing it would measure each of the different reward functions. As a reminder, the Q-learning algorithm was trained with the following reward function:

$$r(s_t, a_t, s_{t+1}) = 3 \cdot n_s + 2 \cdot n_b + n_{hb} - \log_{10}(cyc) \quad (5.1)$$

I would then run the Q-learning algorithm and each of the reward function algorithms measuring their reward during the execution. The decision of whether or not the reward function was correct was then made with the following rule:

If the reward of the deep Q-network algorithm was higher than the one of the Q-learning algorithm, but its overall precision was not higher. This would imply that the deep Q-network was correctly learning to maximise the reward, but that this reward was not maximising the overall performance of the algorithm and therefore that it was not a good reward.

As a reminder, the rewards tested are the following:

$$r_{pr1}(s_t, a_t, s_{t+1}) = 3 \cdot n_s + 2 \cdot n_b + n_{hb} \quad (5.2)$$

$$r_{pr2}(s_t, a_t, s_{t+1}) = 3 \cdot n_s + 2 \cdot n_b + n_{hb} - \log_{10}(|n_b|) \quad (5.3)$$

$$r_{pr3}(s_t, a_t, s_{t+1}) = 3 \cdot (n_{sf} - n_{si}) + 2 \cdot (n_{bf} - n_{bi}) + (n_{hb_f} - n_{hb_i}) - \log_{10}(|n_b|) \quad (5.4)$$

5. Results




Figure 5.1.: Reward one




Figure 5.2.: Reward two




Figure 5.3.: Reward three

These are the graphs of the evolution of the precision reward according to the number of joins. The orange line represents the measured reward whilst running the Q-learning algorithm and the blue line the measured reward of the deep Q-network algorithm trained with this reward as well. The overall precision of the end polyhedra are the following: For Figure 5.1. 99% for Q-learning and 96% for DQN. For figure 5.2. 99% for Q-learning and 94% for DQN. For figure 5.3. 97% for Q-learning and 90% for DQN. Whilst these graphs behave as expected in the first two scenarios, that is that the reward is lower for the algorithm with lower end precision as well. In the last graph we can see something unexpected. Whilst the DQN algorithm has a higher reward throughout most of the execution of the program having an average of 9 versus -8 for the Q-learning algorithm. This tells us that whilst the DQN has learned a correct strategy for maximising the reward, this reward does not guide us to the overall goal of maximising the end precision and therefore reward three is not an optimal reward function. These experiments do not give us much information on the differences between reward one and two as they both behave as normally. This is to be expected as reward one and two are similar.

It is worth noting that whilst these experiments do give us some interesting information about the optimality of the different reward function, the results are heavily dependant of the benchmark we run them on and therefore the results cannot be fully trusted. However, with the information gathered by these experiments as well as some overall testing I decided to use the second reward function in further experiments.

5.2. Action selection algorithm

The second problem was selection the best action selection algorithm. In section 4.8., four different action selection algorithm were described. In order to select the best one of these four, I proceeded by training an algorithm using each of these four selection algorithms and then comparing the overall precision. Not all the training had to be done separately for the four different action selections. Firstly, we could train estimators from only random actions and then simply specialise each of them with their own action selection algorithm. I then compared the results on a set of seven different benchmarks.

Benchmark	Q-learning	Algorithm 1	Algorithm 2	Algorithm 3	Algorithm 4
driver-media	99.0	98.6	97.6	93.2	98.0
linux-kernel-locking-spinlock	99.9	94.4	63.6	79.9	95.0
linux-usb-dev	58.2	51.8	60.3	56.5	59.3
linux-kernel-locking-mutex	77.1	73.2	95.5	95.6	76.1
complex-emg-drivers-net	57.1	97.3	94.6	96.0	96.9
complex-emg-drivers-media	57	50.1	58.9	55.4	54.9
complex-emg-linux-alloc-spinlock	44.9	72.2	69.9	91.8	72.0

As to the choice of benchmarks for this experiment. I had several criteria. First, I picked benchmarks where achieving a high level of accuracy was not too easy in order to have more informative results. Second, I picked benchmarks that were fast to compute so that the testing time would be fast. I also picked fast benchmarks because my goal was to find an algorithm that maximised the precision as each of these algorithms can be then optimised with its own parameters in order to ameliorate performance and in this way I only had to concentrate on precision.

As to the results of these experiments. Unfortunately there is not overall best algorithm. As discussed in section 4.8., each of these algorithms has its advantages and shortcomings. However, the main objective I was looking for from these experiments was finding an algorithm that would be stable. We can see that algorithms two and three can obtain very good on some benchmarks. Unfortunately, they also perform very bad on others. This is due to the fact that both of these are threshold based and picking an optimal threshold that would work well for all benchmarks is very difficult, making both of these algorithms not practical. Afterwards, between the first and fourth algorithm, the last one outperforms the first on average and therefore is the one I will use from here on. Whilst the last algorithm does not have the best precision on some benchmarks its stability over all benchmarks made it the preferred candidate.

5.3. Final algorithm

Once both of these decisions had been made. The architecture of the final separated algorithm was fully decided and could be further trained and optimised.


Conclusion and Future Work

This work was able to design a training algorithm that was capable of dynamically picking transformers with varying precision and performance capabilities. The algorithm was able to learn a decision strategy that is capable of outperforming other such strategies in both precision and performance.

6. *Conclusion and Future Work*

Appendix

in here we put all questionnaires, guides etc.



Bibliography